

Protecting XALT from Users

Robert McLay

April 21, 2022

XALT: Outline



- ▶ XALT is linking with every program that runs on the system
- ▶ Users will occasionally make mistakes
- ▶ Need to protect XALT from user mistakes
- ▶ Show more ways that XALT protects itself and users.

Three examples of protection

- ▶ User's bug hidden by zeroed memory initially
- ▶ User's mixing Fortran routine with C library routines badly
- ▶ XALT expecting well managed memory heap.

How XALT works

```
#include <stdio.h>
void myinit(int argc, char **argv)
{ printf("This is run before main()\n"); }
void myfini()
{ printf("This is run after main()\n"); }

__attribute__((section(".init_array"))) __typeof__(myinit) *__init = myinit;
__attribute__((section(".fini_array"))) __typeof__(myfini) *__fini = myfini;
```

- [my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_memory/ex1](#)

How XALT works (II)

```
% cat try.c
```

```
#include <stdio.h>
int main()
{
    printf("Hello World!\n");
    return 0;
}
```

How XALT works (III)

```
$ ./try
```

Hello World!

```
$ LD_PRELOAD=./libxalt.so ./try
```

This is run before main()

Hello World!

This is run after main()

- ▶ `my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_memory/ex1`

How XALT works (IV)

- ▶ The xalt runtime library attaches to every program run on your system!
- ▶ I sometimes feel like I'm a developer on every program team.
- ▶ XALT programs are in the same namespace as the user's program (UGH!)

Hiding XALT routine names from users

```
% nm $LD_PRELOAD | grep __XALT_build
00000000000009e80 T __XALT_buildEnvT_xalt_1_5
0000000000000a840 T __XALT_buildUserT_xalt_1_5
000000000000163a0 T __XALT_buildXALTRecordT_xalt_1_5
```

- XALT routine names are hidden by macros supplied in `xalt_obfuscate.h`

User's bug hidden by initially zeroed memory

- ▶ Initially all memory is zeroed before program starts
- ▶ Note that pointer zero, integer zero and float zero are all zero bits
- ▶ Link lists require a NULL pointer at end of list.
- ▶ Used memory is *NOT* zeroed for you in C.
- ▶ User's program work w/o XALT, Failed with XALT.

Example code clean/used memory

% cat try.c

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 1000
int main()
{
    int *a = (int *) malloc(SZ*sizeof(int));
    printf("Hello World! a:%d\n",a[0]);
    return 0;
}
```

% cat xalt.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#define SZ 1000
void myinit(int argc, char **argv)
{
    int i;
    int *a = (int*) malloc(SZ*sizeof(int));
    for (i = 0; i < SZ; ++i) a[i] = 15;
    free(a);
    printf("This is run before main()\n");
}
__attribute__((section(".init_array"))) __typeof__(myinit) *__init = myinit;
```

► [my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_memory/ex2](#)

Example code clean/used memory(II)

```
% ./try
```

```
Hello World!  a:0
```

```
% LD_PRELOAD=./libxalt.so  ./try  ; echo
```

```
This is run before main()
```

```
Hello World!  a:15
```

```
This is run after main()
```

- ▶ my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_memory/ex2

XALT Fix: zero memory before free()

- ▶ To protect XALT from broken user code
- ▶ XALT in myinit() zero's memory before free
- ▶ Note that non-MPI tracking does little allocation
- ▶ MPI tasks > 127 init record \Rightarrow much allocation

XALT Fix: zero memory before free()

% cat try.c

```
#include <stdio.h>
#include <stdlib.h>
#define SZ 1000
int main()
{
    int *a = (int *) malloc(SZ*sizeof(int));
    printf("Hello World! a:%d\n",a[0]);
    return 0;
}
```

% cat xalt.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#define SZ 1000
void myinit(int argc, char **argv)
{
    int i;
    int *a = (int*) malloc(SZ*sizeof(int));
    for (i = 0; i < SZ; ++i) a[i] = 15;
    memset((void *) a, 0, SZ*sizeof(int));
    free(a);
    printf("This is run before main()\n");
}
__attribute__((section(".init_array"))) __typeof__(myinit) *__init = myinit;
```

► my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_mem-

XALT Fix: zero memory before free() (II)

```
% ./try
```

```
Hello World!  a:0
```

```
% LD_PRELOAD=./libxalt.so  ./try  ; echo
```

```
This is run before main()
```

```
Hello World!  a:0
```

```
This is run after main()
```

- ▶ `my_docs/22/xalt_monthly_mtg_2022_03_17/code/bad_memory/ex3`

Protecting XALT from Fortran mixed with C programs badly

```
% cat msg.f90
subroutine msg
  print *, "Hello World!"
end subroutine msg
```

```
% nm try | grep msg
000000000000011c7 T msg_
```

- ▶ Normally Fortran routines get a trailing underscore when compiled
- ▶ This can be disabled:
- ▶ Fortran: -fno-underscoring
- ▶ ifort: -assume nounderscore
- ▶ Can make mixing C/Fortran easier
- ▶ Also make collisions with C library easier

XALT uses libuuid

- ▶ libuuid.so is used to get a unique identifier
- ▶ It uses libc's random()
- ▶ Can't have two routines named random()
- ▶ `my_docs/22/xalt_monthly_mtg_2022_03_17/code/random/ex3`

Collision over random() routine

```
% cat try.f90
```

```
program tryMe
  implicit none
  real*8 d, random
  print *, "Hello World!"
  d = random(1.0, 2.0, 3.0)
  print *, "d: ",d
end program tryMe
```

```
% cat random.f90
```

```
real*8 function random(a, b, c)
  implicit none
  real*8 a, b, c
  print *, "In random(a, b, c)"
  random = a*b + c
end function random
```

```
% cat xalt.c
```

```
#include <stdio.h>
#include <stdlib.h>
void myinit(int argc, char **argv)
{
  long int a;
  printf("This is run before main()\n");
  a = random();
  printf("called random(): a: %ld\n",a);
}
__attribute__((section(".init_array"))) __typeof__(myinit) *__init = myinit;
```

Collision over random() routine (II)

```
% ./try

Hello World!
In random(a, b, c)
d:      5.000000000000000000

% LD_PRELOAD=./libxalt.so ./try ; echo
This is run before main()
  In random(a, b, c)
Segmentation fault
```

- ▶ The linker chooses the user's Fortran random() instead of the C lib random()
- ▶ The segfault happens because the Fortran random() expects 3 arguments
- ▶ the random() call in xalt.c passes none.

How to fix this issue

- ▶ Other Fortran program might do the same thing
- ▶ Trick: Use `dlopen()/dlsym()` to dynamically link in `libuuid.so`
- ▶ At this point `libuuid.so` can't "see" the Fortran `random()` routine
- ▶ This trick solves many problems with `libuuid`

XALT is still susceptible to similar issues

- ▶ XALT is now protected from a user's `random()` function
- ▶ But XALT is vulnerable some Fortran code replacing a c library routine
- ▶ We will just have to fix them as they come up

Protecting XALT from badly managed memory heap

```
void my_free(void *ptr,int sz)
{
    if (s_start_record && ptr != NULL)
    {
        memset(ptr, '\0', sz);
        free(ptr);
    }
}
```

- ▶ Reporting an end record in myfini() requires memory allocations
- ▶ However some user programs can leave the heap broken
- ▶ XALT replaces free() with my_free()
- ▶ Memory is only freed for a start record.

Containers

- ▶ Containers are chrooted environments
- ▶ They can have minimum setup and libraries
- ▶ Typically only what the user needs
- ▶ But not always what XALT needs.
- ▶ The xalt library needs many libraries libcrypto, libuuid, libcurl, etc.

Containers (II)

- ▶ For XALT to work at all in a container. LD_PRELOAD must be set.
- ▶ And add XALT_INSTALL_DIR to container's available paths.
- ▶ So XALT during installation copies system libraries to the XALT_INSTALL_DIR
- ▶ XALT does a ldd on executables and libraries to find all the system libraries that need to be copied.
- ▶ This allows XALT to dlopen()/dlsym() libuuid.so to solve the problems mentioned before.

MPI libraries disallow system() after main() returns

- ▶ One of the version of IMPI, stopped supporting system() being called after main() finishes
- ▶ This was a disaster for collecting data from MPI programs.
- ▶ It even prevented collecting data from one task mpi programs
- ▶ This was a problem because XALT used C++ programs to collect data.

Disallowed system() after main()

- ▶ XALT needs Hash Table (AKA Dictionaries) to store many key-value pairs.
- ▶ C++ has built-in hash tables via the STL (unordered-maps)
- ▶ C does not.
- ▶ The XALT library (libxalt_initialize.so) is written in C
- ▶ It would be a bad idea to have written the library in C++

Solution

- ▶ Re-write the XALT submission programs into C
- ▶ And include them in libxalt_initialize.so
- ▶ Where to find a hash table implementation
- ▶ uthash

UTHASH and friends

- ▶ UTHASH: github: <https://github.com/troydhanson/uthash>
- ▶ UTHASH: Docs:
<https://troydhanson.github.io/uthash/userguide.html>
- ▶ uthash is written in C via uthash.h (no library!)

UTHASH

```
// One C++ stmt:
userDT["start_time"] = start_time;

// Becomes several C stmts:
insert_key_double(&userDT, "start_time", start_time);
// -->
void insert_key_double(S2D_t** userDT, const char* name, double value)

    S2D_t* entry = (S2D_t *) XMALLOC(sizeof(S2D_t));
    utstring_new(entry->key);
    utstring_bincpy(entry->key, name, strlen(name));

    entry->value = value;

    HASH_ADD_KEYPTR(hh, *userDT, utstring_body(entry->key),
                    utstring_len(entry->key), entry);
```

UTHASH rewrite

- ▶ It took a couple weeks to rewrite the submission programs into routines
- ▶ The only system calls that still exist are the curl calls for transport.

Support for filtering of packages

- ▶ Python packages could be directly filtered through the `sitecustomize.py` file
- ▶ But not R or MATLAB packages
- ▶ All packages collection goes through the `xalt_record_pkg` program

All package filtering

- ▶ New pkg_pattern array available in XALT 2.10.37+
- ▶ This goes in your site Config.py
- ▶ The pattern looks like this:
- ▶ Program:kind:pattern

```
pkg_patterns = [  
    ["SKIP", r'^R:name:stats'],          # SKIP the R pkg named stats  
    ["SKIP", r'^R:name:base'],           # SKIP the R pkg named base  
    ["SKIP", r'^R:name:methods'],        # SKIP the R pkg named methods  
    ["SKIP", r'^python:name:_.*'],       # SKIP all python name that start with an '_'  
    ["SKIP", r'^python:path:[^/].*'],    # SKIP all python built-in packages  
    ["SKIP", r'^python:path:\home'],     # SKIP all python package in user locations  
]
```

This package filtering could be added to ingestion

- ▶ The same routine could be added to ingestion.
- ▶ If wanted it could be added in late May 2022.

Conclusions

- ▶ XALT has matured greatly from working with user programs
- ▶ Since the XALT library is in the same namespace as the user code
- ▶ There is always a risk of routine collision.

Future Topics?

- ▶ Recent changes to importing json records
- ▶ Others?